

Time Series Analysis - Time_Series_Analysis_Daily_Website_Visitors

Introduction to Time-Series Analysis

- A time-series data is a series of data points or observations recorded at different or regular time intervals. In general, a time series is a sequence of data points taken at equally spaced time intervals. The frequency of recorded data points may be hourly, daily, weekly, monthly, quarterly or annually.
- Time-Series Forecasting is the process of using a statistical model to predict future values of a time-series based on past results. • A time series analysis encompasses statistical methods for analyzing time series data. These methods enable us to extract meaningful statistics, patterns and other characteristics of the data. Time series are visualized with the help of line charts. So, time series analysis involves understanding inherent aspects of the time series data so that we can create meaningful and accurate forecasts.
- Applications of time series are used in statistics, finance or business applications. A very common example of time series data is the daily closing value of the stock index like NASDAQ or Dow Jones. Other common applications of time series are sales and demand forecasting, weather forecasting, econometrics, signal processing, pattern recognition and earthquake prediction

Components of a Time-Series

- Trend - The trend shows a general direction of the time series data over a long period of time. A trend can be increasing(upward), decreasing(downward), or horizontal(stationary).
- Seasonality - The seasonality component exhibits a trend that repeats with respect to timing, direction, and magnitude. Some examples include an increase in water consumption in summer due to hot weather conditions.
- Cyclical Component - These are the trends with no set repetition over a particular period of time. A cycle refers to the period of ups and downs, booms and slumps of a time series, mostly observed in business cycles. These cycles do not exhibit a seasonal variation but generally occur over a time period of 3 to 12 years depending on the nature of the time series.
- Irregular Variation - These are the fluctuations in the time series data which become evident when trend and cyclical variations are removed. These variations are unpredictable, erratic, and may or may not be random.
- ETS Decomposition - ETS Decomposition is used to separate different components of a time series. The term ETS stands for Error, Trend and Seasonality

```
In [ ]: ##!mkdir ~/.kaggle
```

```
In [4]: ##!cp /kaggle.json ~/.kaggle/
```

```
In [5]: ##!chmod 600 ~/.kaggle/kaggle.json
```

```
In [ ]: ####! pip install kaggle
```

```
In [ ]: ####!pip install keras-tuner
```

```
In [ ]: ####! kaggle datasets download -d bobnau/daily-website-visitors
```

```
In [10]: ####! unzip ./daily-website-visitors.zip
```

```
In [11]: ####! pip install tensorflow
```

```
In [12]: import pandas as pd  
import numpy as np  
import tensorflow as tf  
from sklearn import preprocessing  
import matplotlib.pyplot as plt  
tf.random.set_seed(123)  
np.random.seed(123)
```

```
In [13]: import pandas as pd  
         #import fbprophet  
import matplotlib.pyplot as plt  
         %matplotlib inline  
import numpy as np
```

```
In [14]: df=pd.read_csv('/content/daily-website-visitors.csv', parse_dates=['Date'], index_col = "Date")
df.head(3)
```

```
Out[14]:
```

	Row	Day	Day.Of.Week	Page.Loads	Unique.Visits	First.Time.Visits	Returning.Visits	
	Date							
	2014-09-14	1	Sunday	1	2,146	1,582	1,430	152
	2014-09-15	2	Monday	2	3,621	2,528	2,297	231
	2014-09-16	3	Tuesday	3	3,698	2,630	2,352	278

```
In [15]: df.columns
```

```
Out[15]: Index(['Row', 'Day', 'Day.Of.Week', 'Page.Loads', 'Unique.Visits',
               'First.Time.Visits', 'Returning.Visits'],
              dtype='object')
```

```
In [16]: df2 = df [['Page.Loads', 'Unique.Visits',
                   'First.Time.Visits', 'Returning.Visits']]
```

```
In [17]: df2.describe()
```

```
Out[17]:
```

	Page.Loads	Unique.Visits	First.Time.Visits	Returning.Visits
count	2167	2167	2167	2167
unique	1756	1658	1587	663
top	2,948	2,780	3,146	552
freq	5	5	5	12

```
In [18]: df2['Page.Loads'] = df2['Page.Loads'].str.replace(',', '').astype(int)
```

```
<ipython-input-18-10a780dcfd9e>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df2['Page.Loads'] = df2['Page.Loads'].str.replace(',', ' ').astype(int)
```

```
In [19]: df2['Unique.Visits'] = df2['Unique.Visits'].str.replace(',', ' ').astype(int)
```

```
<ipython-input-19-749f4711fe3b>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df2['Unique.Visits'] = df2['Unique.Visits'].str.replace(',', ' ').astype(int)
```

```
In [20]: df2['First.Time.Visits'] = df2['First.Time.Visits'].str.replace(',', ' ').astype(int)
```

```
<ipython-input-20-d715d5f15c0d>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df2['First.Time.Visits'] = df2['First.Time.Visits'].str.replace(',', ' ').astype(int)
```

```
In [21]: df2['Returning.Visits'] = df2['Returning.Visits'].str.replace(',', ' ').astype(int)
```

```
<ipython-input-21-7e6c0f8bfde5>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
df2['Returning.Visits'] = df2['Returning.Visits'].str.replace(',', ' ').astype(int)
```

```
In [22]: df2.head(3)
```

Out[22]:

	Page.Loads	Unique.Visits	First.Time.Visits	Returning.Visits
Date				
2014-09-14	2146	1582	1430	152
2014-09-15	3621	2528	2297	231
2014-09-16	3698	2630	2352	278

Date				
2014-09-14	2146	1582	1430	152
2014-09-15	3621	2528	2297	231
2014-09-16	3698	2630	2352	278

```
In [23]: # Basic packages
import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import random as rd # generating random numbers
import datetime # manipulating date formats
# Viz
import matplotlib.pyplot as plt # basic plotting
import seaborn as sns # for prettier plots
# TIME SERIES
from statsmodels.tsa.arima_model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from pandas.plotting import autocorrelation_plot
from statsmodels.tsa.stattools import adfuller, acf, pacf, arma_order_select_ic
import statsmodels.formula.api as smf
import statsmodels.tsa.api as smt
import statsmodels.api as sm
import scipy.stats as scs
# settings
import warnings
warnings.filterwarnings("ignore")
```

```
In [24]: monthly_page_loads = df2["Page.Loads"].resample('M').sum()
```

```
In [25]: monthly_unique_visitors = df2["Unique.Visits"].resample('M').sum()
```

```
In [26]: monthly_First_Time_Visits = df2["First.Time.Visits"].resample('M').sum()
```

```
In [27]: monthly_Returning_Visits = df2["Returning.Visits"].resample('M').sum()
```

```
In [28]: monthly_page_loads.head(3)
```

```
Out[28]: Date
2014-09-30    56052
2014-10-31    121983
2014-11-30    114190
Freq: M, Name: Page.Loads, dtype: int64
```

```
In [29]: monthly_page_loads = pd.DataFrame(monthly_page_loads)
```

```
In [30]: monthly_page_loads.head(2)
```

```
Out[30]:
```

	Date	Page.Loads
	2014-09-30	56052
	2014-10-31	121983

```
In [31]: monthly_page_loads = monthly_page_loads.reset_index()
```

```
In [32]: monthly_page_loads['weekday'] = monthly_page_loads['Date'].apply(lambda x: x.weekday())
monthly_page_loads.head()
```

```
Out[32]:
```

	Date	Page.Loads	weekday
0	2014-09-30	56052	1
1	2014-10-31	121983	4
2	2014-11-30	114190	6
3	2014-12-31	105617	2
4	2015-01-31	96077	5

```
In [33]: monthly_page_loads['month'] = monthly_page_loads['Date'].dt.month
monthly_page_loads.head()
```

Out[33]:

	Date	Page.Loads	weekday	month
0	2014-09-30	56052	1	9
1	2014-10-31	121983	4	10
2	2014-11-30	114190	6	11
3	2014-12-31	105617	2	12
4	2015-01-31	96077	5	1

```
In [34]: monthly_page_loads['day']=monthly_page_loads['Date'].dt.day  
monthly_page_loads.head()
```

Out[34]:

	Date	Page.Loads	weekday	month	day
0	2014-09-30	56052	1	9	30
1	2014-10-31	121983	4	10	31
2	2014-11-30	114190	6	11	30
3	2014-12-31	105617	2	12	31
4	2015-01-31	96077	5	1	31

```
In [35]: train_month = monthly_page_loads.groupby(["month", "weekday"])['Page.Loads'].mean().reset_index()  
train_month = train_month.pivot('weekday', 'month', 'Page.Loads')  
train_month.sort_index(inplace=True)
```


```
In [36]: train_month.head()
```

```
Out[36]:
```

month	1	2	3	4	5	6	7	8	9	10	11	12
weekday												
0	NaN	131171.0	NaN	175572.0	NaN	NaN	73308.0	79732.0	123213.0	157604.0	142650.0	139140.0
1	118114.0	122170.0	132481.5	167945.0	151279.0	119990.0	101669.0	NaN	56052.0	102451.0	NaN	124419.0
2	111716.0	121740.0	NaN	NaN	119004.0	NaN	103522.0	114891.0	114716.0	171488.0	159651.0	105617.0
3	129723.0	139744.0	145948.0	152386.5	161708.0	125190.0	NaN	72854.0	NaN	148176.0	112155.0	122397.0
4	106492.0	NaN	125290.0	NaN	140413.0	87766.0	104132.0	109901.0	126234.0	121983.0	167131.0	NaN


```
In [37]: import seaborn as sns
```

```
sns.set(font_scale=3.5)
plt.style.use('dark_background')
# Draw a heatmap with the numeric values in each cell
f, ax = plt.subplots(figsize=(50, 15))
sns.heatmap(train_month, annot=False, ax=ax, fmt="d", linewidths=10)
plt.title('Confirmation Of Months cross Weekdays')
plt.show()
```




```
In [38]: plt.figure(figsize=(50,15))  
plt.style.use('dark_background')  
sns.violinplot(df2['Page.Loads'])
```

Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2cb34c1c10>



```
In [39]: monthly_page_loads.head(10)
```

Out[39]:

	Date	Page.Loads	weekday	month	day
0	2014-09-30	56052	1	9	30
1	2014-10-31	121983	4	10	31
2	2014-11-30	114190	6	11	30
3	2014-12-31	105617	2	12	31
4	2015-01-31	96077	5	1	31
5	2015-02-28	118876	5	2	28
6	2015-03-31	143990	1	3	31
7	2015-04-30	153331	3	4	30
8	2015-05-31	142113	6	5	31
9	2015-06-30	115297	1	6	30

```
In [40]: train_days = monthly_page_loads.groupby(["month", "day"])['Page.Loads'].mean().reset_index()
train_days = train_days.pivot('day', 'month', 'Page.Loads')
train_days.sort_index(inplace=True)
train_days.dropna(inplace=True)
```

```
In [41]: df2.columns
```

```
Out[41]: Index(['Page.Loads', 'Unique.Visits', 'First.Time.Visits', 'Returning.Visits'], dtype='object')
```

```
In [42]: import statsmodels.api as sm
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
In [43]: #####! pip install plotly
```

```
In [44]: import plotly
# plotly.tools.set_credentials_file()
```

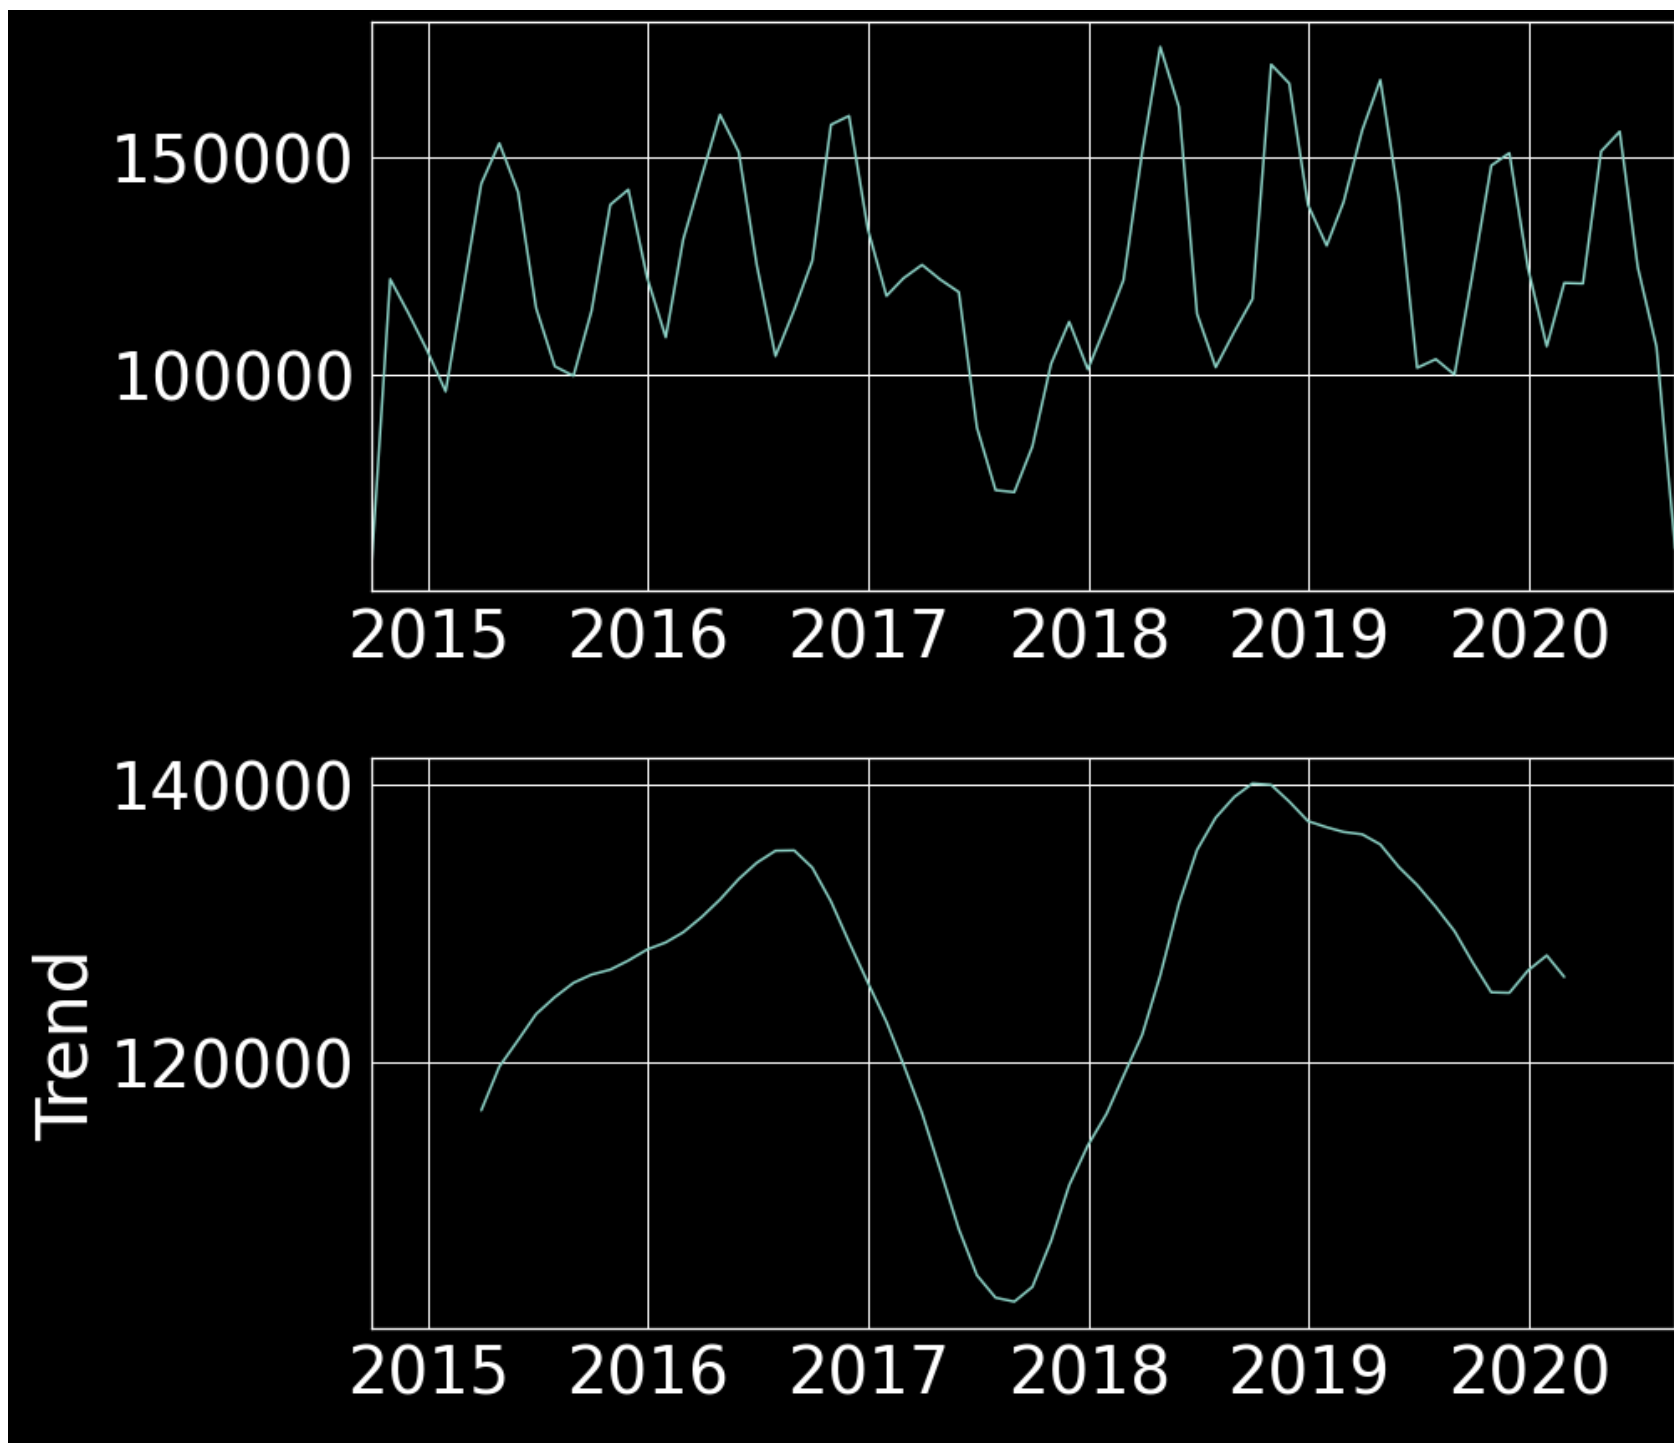
```
In [45]: monthly_page_loads.rename(columns = {"Page.Loads" : "Page_Loads"}, inplace=True)
```

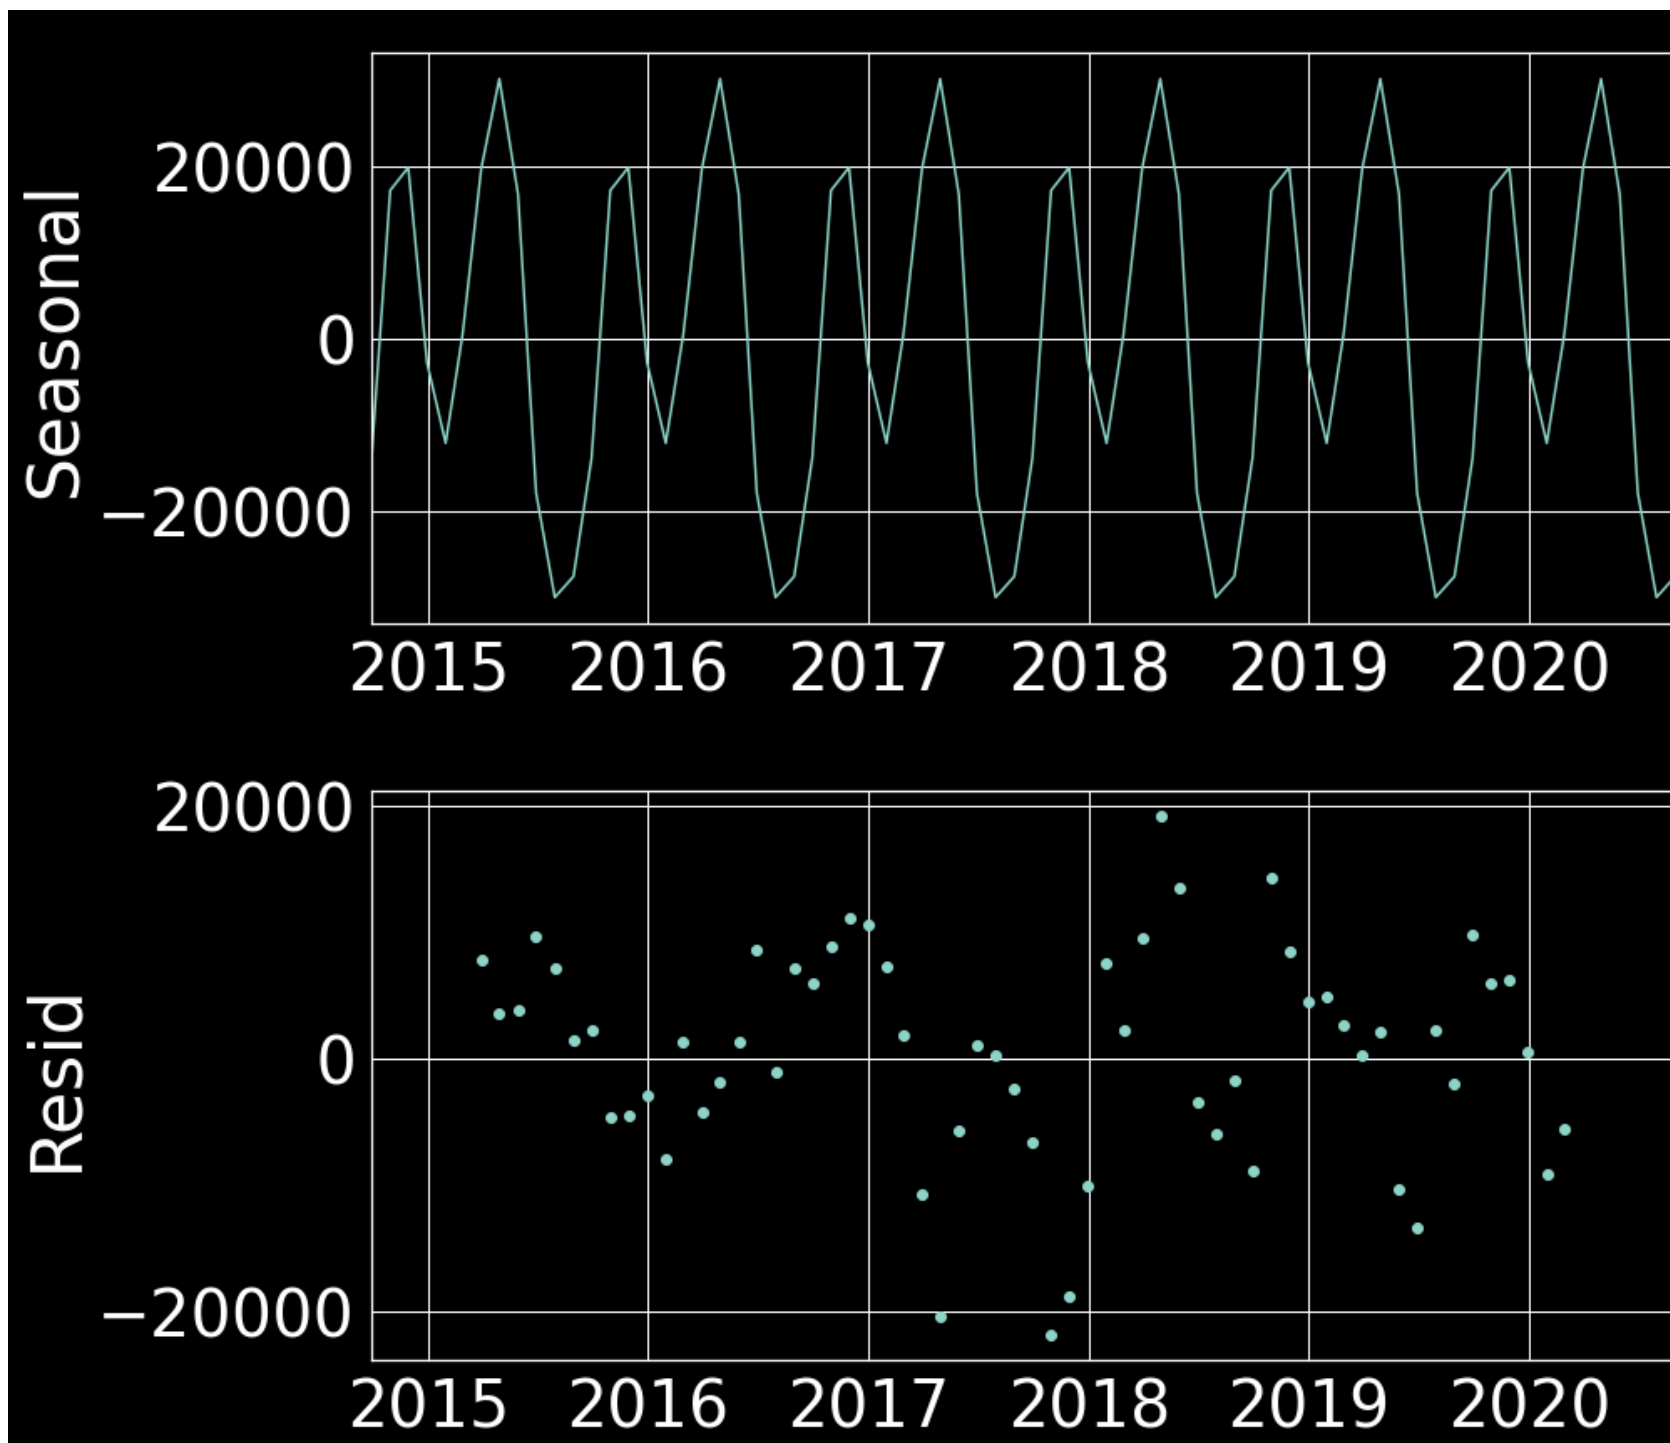
```
In [46]: # Show Rolling mean, Rolling Std and Test for the stationnarity
df_date_index = monthly_page_loads[['Date', 'Page_Loads']].set_index('Date')
df_date_index.head()
```

```
Out[46]:
```

	Page_Loads
Date	
2014-09-30	56052
2014-10-31	121983
2014-11-30	114190
2014-12-31	105617
2015-01-31	96077

```
In [47]: from pylab import rcParams
rcParams['figure.figsize'] = 15, 25
decomposition = sm.tsa.seasonal_decompose(df_date_index, model='additive')
fig = decomposition.plot()
plt.show()
```





Stationarity


A Time Series is said to be stationary if its statistical properties such as mean, variance remain constant over time. Most of the Time Series models work on the assumption that the TS is stationary. Major reason for this is that there are many ways in which a series can be non-stationary, but only one way for stationarity.

Intuitively, we can say that if a Time Series has a particular behaviour over time, there is a very high probability that it will follow the same in the future.

Also, the theories related to stationary series are more mature and easier to implement as compared to non-stationary series.

```
In [48]: import matplotlib.pyplot as plt  
plt.rcParams["figure.figsize"] = (20,3)  
plt.plot(df_date_index)
```

```
Out[48]: [<matplotlib.lines.Line2D at 0x7f2cacdf0f40>]
```



Testing For Stationarity

```
In [49]: ### Testing For Stationarity

from statsmodels.tsa.stattools import adfuller

def adf_test(dataset):
    dfctest = adfuller(dataset, autolag = 'AIC')
    print("1. ADF : ",dfctest[0])
    print("2. P-Value : ", dfctest[1])
    print("3. Num Of Lags : ", dfctest[2])
    print("4. Num Of Observations Used For ADF Regression and Critical Values Calculation :", dfctest[3])
    print("5. Critical Values :")
    for key, val in dfctest[4].items():
        print("\t",key, ": ", val)
```

AD-FULLER-TEST

```
In [50]: df_date_index.columns
```

```
Out[50]: Index(['Page_Loads'], dtype='object')
```

```
In [51]: adf_test(df_date_index['Page_Loads'])
```

```
1. ADF : -2.410024672338927
2. P-Value : 0.13892698645568835
3. Num Of Lags : 12
4. Num Of Observations Used For ADF Regression and Critical Values Calculation : 59
5. Critical Values :
    1% : -3.5463945337644063
    5% : -2.911939409384601
   10% : -2.5936515282964665
```

DIFFERENCING

```
In [52]: df_date_index['Page_Loads_First_Order_Differencing'] = df_date_index['Page_Loads'] - df_date_index['Page_Loads'].shift(1)
```

```
In [53]: adf_test(df_date_index['Page_Loads_First_Order_Differencing'].dropna())
```

```
1. ADF : -1.9067763211534854
2. P-Value : 0.32888006117831814
3. Num Of Lags : 11
4. Num Of Observations Used For ADF Regression and Critical Values Calculation : 59
5. Critical Values :
    1% : -3.5463945337644063
    5% : -2.911939409384601
    10% : -2.5936515282964665
```

```
In [54]: df_date_index['Page_Loads_Second_Order_Differencing'] = df_date_index['Page_Loads_First_Order_Differencing'] - df_date_index['Page_Loads_Second_Order_Differencing']
```

```
In [55]: adf_test(df_date_index['Page_Loads_Second_Order_Differencing'].dropna())
```

```
1. ADF : -5.5736290225934155
2. P-Value : 1.4486118892546306e-06
3. Num Of Lags : 10
4. Num Of Observations Used For ADF Regression and Critical Values Calculation : 59
5. Critical Values :
    1% : -3.5463945337644063
    5% : -2.911939409384601
    10% : -2.5936515282964665
```

```
In [56]: def adfuller_test(confirmed):

    result=adfuller(confirmed)

    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of Observations Used']

    for value,label in zip(result,labels):
        print(label+' : '+str(value) )

    if result[1] <= 0.05:

        print("strong evidence against the null hypothesis(Ho), reject the null hypothesis. Data has no unit root and is stationary")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary ")
```



```
In [57]: adfuller_test(df_date_index['Page_Loads_Second_Order_Differencing'].dropna())
```

ADF Test Statistic : -5.5736290225934155

p-value : 1.4486118892546306e-06

#Lags Used : 10

Number of Observations Used : 59

strong evidence against the null hypothesis(H_0), reject the null hypothesis. Data has no unit root and is stationary

Plotting ACF and PACF

Autocorrelation and Partial Autocorrelation Functions


- Autocorrelation is simply the correlation of a series with its own lags. If a series is significantly autocorrelated, that means, the previous values of the series (lags) may be helpful in predicting the current value.
- Partial Autocorrelation also conveys similar information but it conveys the pure correlation of a series and its lag, excluding the correlation contributions from the intermediate lags.

```
In [58]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```


```
In [59]: df_date_index.columns
```

```
Out[59]: Index(['Page_Loads', 'Page_Loads_First_Order_Differencing',  
              'Page_Loads_Second_Order_Differencing'],  
             dtype='object')
```


```
In [60]: acf2 = plot_acf(df_date_index['Page_Loads_Second_Order_Differencing'].dropna())
```




```
In [61]: acf1 = plot_acf(df_date_index["Page_Loads_First_Order_Differencing"].dropna())
```



```
In [62]: result1 = plot_pacf(df_date_index["Page_Loads_First_Order_Differencing"].dropna())
```



```
In [63]: result2 = plot_pacf(df_date_index['Page_Loads_Second_Order_Differencing'].dropna())
```



Split Data into Training and Testing

```
In [64]: print(df_date_index.shape)
train=df_date_index.iloc[:-30]
test=df_date_index.iloc[-30:]
print(train.shape,test.shape)
### print(test.iloc[0],test.iloc[-1])
```

```
(72, 3)
(42, 3) (30, 3)
```

```
In [65]: train.columns
```

```
Out[65]: Index(['Page_Loads', 'Page_Loads_First_Order_Differencing',
               'Page_Loads_Second_Order_Differencing'],
              dtype='object')
```

```
In [66]: from statsmodels.tsa.arima.model import ARIMA
model_ARIMA=ARIMA(train['Page_Loads_Second_Order_Differencing'],order=(0,2,0))
```

```
In [67]: model_Arima_fit=model_ARIMA.fit()
```

```
In [68]: model_Arima_fit.summary()
```

Out[68]:

SARIMAX Results

Dep. Variable:	Page_Loads_Second_Order_Differencing	No. Observations:	42
Model:	ARIMA(0, 2, 0)	Log Likelihood	-483.086
Date:	Sun, 11 Dec 2022	AIC	968.171
Time:	16:08:09	BIC	969.860
Sample:	09-30-2014	HQIC	968.782
	- 02-28-2018		
Covariance Type:	opg		
	coef	std err	z P> z [0.025 0.975]
sigma2	1.381e+09	3.09e+08	4.465 0.000 7.75e+08 1.99e+09
Ljung-Box (L1) (Q):	13.79	Jarque-Bera (JB):	0.70
Prob(Q):	0.00	Prob(JB):	0.70
Heteroskedasticity (H):	0.87	Skew:	-0.06
Prob(H) (two-sided):	0.80	Kurtosis:	2.36

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [69]: ##prediction
pred_start_date=test.index[0]
pred_end_date=test.index[-1]
print(pred_start_date)
print(pred_end_date)
```

```
2018-03-31 00:00:00
2020-08-31 00:00:00
```


```
In [70]: pred=model_Arima_fit.predict(start=pred_start_date,end=pred_end_date)
```

```
In [71]: residuals=test['Page_Loads']- pred
```

```
In [72]: residuals.head(3)
```

```
Out[72]: Date
2018-03-31    173775.0
2018-04-30    220193.0
2018-05-31    228379.0
dtype: float64
```

```
In [73]: plt.figure(figsize=[7,5])
model_Arima_fit.resid.plot(kind='kde')
plt.show()
```



```
In [74]: test['Predicted_ARIMA']=pred
```

```
In [75]: test.columns
```

```
Out[75]: Index(['Page_Loads', 'Page_Loads_First_Order_Differencing',
              'Page_Loads_Second_Order_Differencing', 'Predicted_ARIMA'],
              dtype='object')
```

```
In [79]: from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, median_absolute_error, mean_squared_lo
```

```
In [80]: from math import sqrt
```


```
In [81]: # report performance
mse = mean_squared_error(test["Page_Loads"], test["Predicted_ARIMA"])
rmse = sqrt(mse)
print('ARIMA RMSE: {}, MSE:{}'.format(rmse,mse))

##plt.title('RMSE: %.4f'% rmse)
```

```
ARIMA RMSE: 507774.5966582679, MSE:257835041011.46667
```

```
In [82]: start=len(train)
end=len(train)+len(test)-1
#if the predicted values dont have date values as index, you will have to uncomment the following two commented lines
#index_future_dates=pd.date_range(start='2018-12-01',end='2018-12-30')
pred=model_Arima_fit.predict(start=start,end=end,typ='levels').rename('ARIMA predictions')
#pred.index=index_future_dates
pred.plot(legend=True)
test['Page_Loads'].plot(legend=True)
```

```
Out[82]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2cabdd1610>
```



```
In [83]: df2.columns
```

```
Out[83]: Index(['Page.Loads', 'Unique.Visits', 'First.Time.Visits', 'Returning.Visits'], dtype='object')
```

```
In [84]: df2.head(3)
```

```
Out[84]:
```

	Page.Loads	Unique.Visits	First.Time.Visits	Returning.Visits
Date				
2014-09-14	2146	1582	1430	152
2014-09-15	3621	2528	2297	231
2014-09-16	3698	2630	2352	278

```
In [85]: df2 = df2.reset_index()
```

```
In [86]: df2.head(3)
```

```
Out[86]:
```

	Date	Page.Loads	Unique.Visits	First.Time.Visits	Returning.Visits
0	2014-09-14	2146	1582	1430	152
1	2014-09-15	3621	2528	2297	231
2	2014-09-16	3698	2630	2352	278

```
In [87]: final = df2[["Date", "Page.Loads"]]
```

```
In [88]: final.columns
```

```
Out[88]: Index(['Date', 'Page.Loads'], dtype='object')
```

```
In [89]: y = pd.Series(data=final['Page.Loads'].values, index=final['Date'])
```

```
In [90]: y.head(3)
```

```
Out[90]: Date
2014-09-14    2146
2014-09-15    3621
2014-09-16    3698
dtype: int64
```

Grid search the `p`, `d`, `q` parameters

```
In [91]: import itertools
# Define the p, d and q parameters to take any value between 0 and 3
p = d = q = range(0, 2)

# Generate all different combinations of p, q and q triplets
pdq = list(itertools.product(p, d, q))

# Generate all different combinations of seasonal p, q and q triplets
seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
```



```
In [92]: warnings.filterwarnings("ignore") # specify to ignore warning messages

best_result = [0, 0, 1000]
for param in pdq:
    for param_seasonal in seasonal_pdq:
        try:
            mod = sm.tsa.statespace.SARIMAX(y,
                                             order=param,
                                             seasonal_order=param_seasonal,
                                             enforce_stationarity=False,
                                             enforce_invertibility=False)

            results = mod.fit()

            print('ARIMA{} x {} - AIC: {}'.format(param, param_seasonal, results.aic))

            if results.aic < best_result[2]:
                best_result = [param, param_seasonal, results.aic]
        except:
            continue

print('\nBest Result:', best_result)
```

ARIMA(0, 0, 0) x (0, 0, 0, 12) - AIC: 42425.717184898036
ARIMA(0, 0, 0) x (0, 0, 1, 12) - AIC: 40851.64824104455
ARIMA(0, 0, 0) x (0, 1, 0, 12) - AIC: 38145.78754098146
ARIMA(0, 0, 0) x (0, 1, 1, 12) - AIC: 37044.05664493689
ARIMA(0, 0, 0) x (1, 0, 0, 12) - AIC: 38080.43136062496
ARIMA(0, 0, 0) x (1, 0, 1, 12) - AIC: 37277.31040369803
ARIMA(0, 0, 0) x (1, 1, 0, 12) - AIC: 37749.32223654218
ARIMA(0, 0, 0) x (1, 1, 1, 12) - AIC: 36956.07859514652
ARIMA(0, 0, 1) x (0, 0, 0, 12) - AIC: 40710.94920795676
ARIMA(0, 0, 1) x (0, 0, 1, 12) - AIC: 37771.581807463284
ARIMA(0, 0, 1) x (0, 1, 0, 12) - AIC: 36401.431809061185
ARIMA(0, 0, 1) x (0, 1, 1, 12) - AIC: 35111.40389091437
ARIMA(0, 0, 1) x (1, 0, 0, 12) - AIC: 36250.642693204965
ARIMA(0, 0, 1) x (1, 0, 1, 12) - AIC: 35391.14440267109
ARIMA(0, 0, 1) x (1, 1, 0, 12) - AIC: 35884.466502963856
ARIMA(0, 0, 1) x (1, 1, 1, 12) - AIC: 35066.32474097704
ARIMA(0, 1, 0) x (0, 0, 0, 12) - AIC: 35890.78174577265
ARIMA(0, 1, 0) x (0, 0, 1, 12) - AIC: 34885.16639554426
ARIMA(0, 1, 0) x (0, 1, 0, 12) - AIC: 37860.90404717403
ARIMA(0, 1, 0) x (0, 1, 1, 12) - AIC: 36294.68822314573
ARIMA(0, 1, 0) x (1, 0, 0, 12) - AIC: 35402.969076336805
ARIMA(0, 1, 0) x (1, 0, 1, 12) - AIC: 34884.424872536256
ARIMA(0, 1, 0) x (1, 1, 0, 12) - AIC: 37051.814098113595
ARIMA(0, 1, 0) x (1, 1, 1, 12) - AIC: 35249.199104267274
ARIMA(0, 1, 1) x (0, 0, 0, 12) - AIC: 35304.602680482
ARIMA(0, 1, 1) x (0, 0, 1, 12) - AIC: 34670.40680000931
ARIMA(0, 1, 1) x (0, 1, 0, 12) - AIC: 37118.317716272315
ARIMA(0, 1, 1) x (0, 1, 1, 12) - AIC: 34972.45745761369
ARIMA(0, 1, 1) x (1, 0, 0, 12) - AIC: 34910.486235664896
ARIMA(0, 1, 1) x (1, 0, 1, 12) - AIC: 34672.38508734617
ARIMA(0, 1, 1) x (1, 1, 0, 12) - AIC: 36157.23287649274
ARIMA(0, 1, 1) x (1, 1, 1, 12) - AIC: 34740.74008456864
ARIMA(1, 0, 0) x (0, 0, 0, 12) - AIC: 35883.807465184334
ARIMA(1, 0, 0) x (0, 0, 1, 12) - AIC: 34902.16793390615
ARIMA(1, 0, 0) x (0, 1, 0, 12) - AIC: 37345.03899148312
ARIMA(1, 0, 0) x (0, 1, 1, 12) - AIC: 35288.052190881994
ARIMA(1, 0, 0) x (1, 0, 0, 12) - AIC: 35411.4277048134
ARIMA(1, 0, 0) x (1, 0, 1, 12) - AIC: 35029.9320734105
ARIMA(1, 0, 0) x (1, 1, 0, 12) - AIC: 36639.60274432837
ARIMA(1, 0, 0) x (1, 1, 1, 12) - AIC: 35078.57653277343
ARIMA(1, 0, 1) x (0, 0, 0, 12) - AIC: 35272.22925242842
ARIMA(1, 0, 1) x (0, 0, 1, 12) - AIC: 34745.1460108378
ARIMA(1, 0, 1) x (0, 1, 0, 12) - AIC: 36229.57075185914
ARIMA(1, 0, 1) x (0, 1, 1, 12) - AIC: 34484.598122057796

```

ARIMA(1, 0, 1) x (1, 0, 0, 12) - AIC: 34888.19402088795
ARIMA(1, 0, 1) x (1, 0, 1, 12) - AIC: 34747.13227962806
ARIMA(1, 0, 1) x (1, 1, 0, 12) - AIC: 35519.09094533223
ARIMA(1, 0, 1) x (1, 1, 1, 12) - AIC: 34412.95268709655
ARIMA(1, 1, 0) x (0, 0, 0, 12) - AIC: 35708.171394599536
ARIMA(1, 1, 0) x (0, 0, 1, 12) - AIC: 34787.27113789217
ARIMA(1, 1, 0) x (0, 1, 0, 12) - AIC: 37738.08264166268
ARIMA(1, 1, 0) x (0, 1, 1, 12) - AIC: 35373.812247399575
ARIMA(1, 1, 0) x (1, 0, 0, 12) - AIC: 35115.198460961175
ARIMA(1, 1, 0) x (1, 0, 1, 12) - AIC: 34785.742504817164
ARIMA(1, 1, 0) x (1, 1, 0, 12) - AIC: 36627.89952571115
ARIMA(1, 1, 0) x (1, 1, 1, 12) - AIC: 34976.89381286328
ARIMA(1, 1, 1) x (0, 0, 0, 12) - AIC: 35276.82667338982
ARIMA(1, 1, 1) x (0, 0, 1, 12) - AIC: 34651.79421762709
ARIMA(1, 1, 1) x (0, 1, 0, 12) - AIC: 37067.90551756926
ARIMA(1, 1, 1) x (0, 1, 1, 12) - AIC: 34945.14396482198
ARIMA(1, 1, 1) x (1, 0, 0, 12) - AIC: 34891.79033229842
ARIMA(1, 1, 1) x (1, 0, 1, 12) - AIC: 34652.11602235254
ARIMA(1, 1, 1) x (1, 1, 0, 12) - AIC: 36140.28795494583
ARIMA(1, 1, 1) x (1, 1, 1, 12) - AIC: 34738.81469452181

```

Best Result: [0, 0, 1000]

In [93]: `train.head(3)`

Out[93]:

	Page_Loads	Page_Loads_First_Order_Differencing	Page_Loads_Second_Order_Differencing
Date			
2014-09-30	56052	NaN	NaN
2014-10-31	121983	65931.0	NaN
2014-11-30	114190	-7793.0	-73724.0

SARIMAX

The implementation is called SARIMAX instead of SARIMA because the "X" addition to the method name means that the implementation also supports exogenous variables. Exogenous variables are optional can be specified via the "exog" argument.

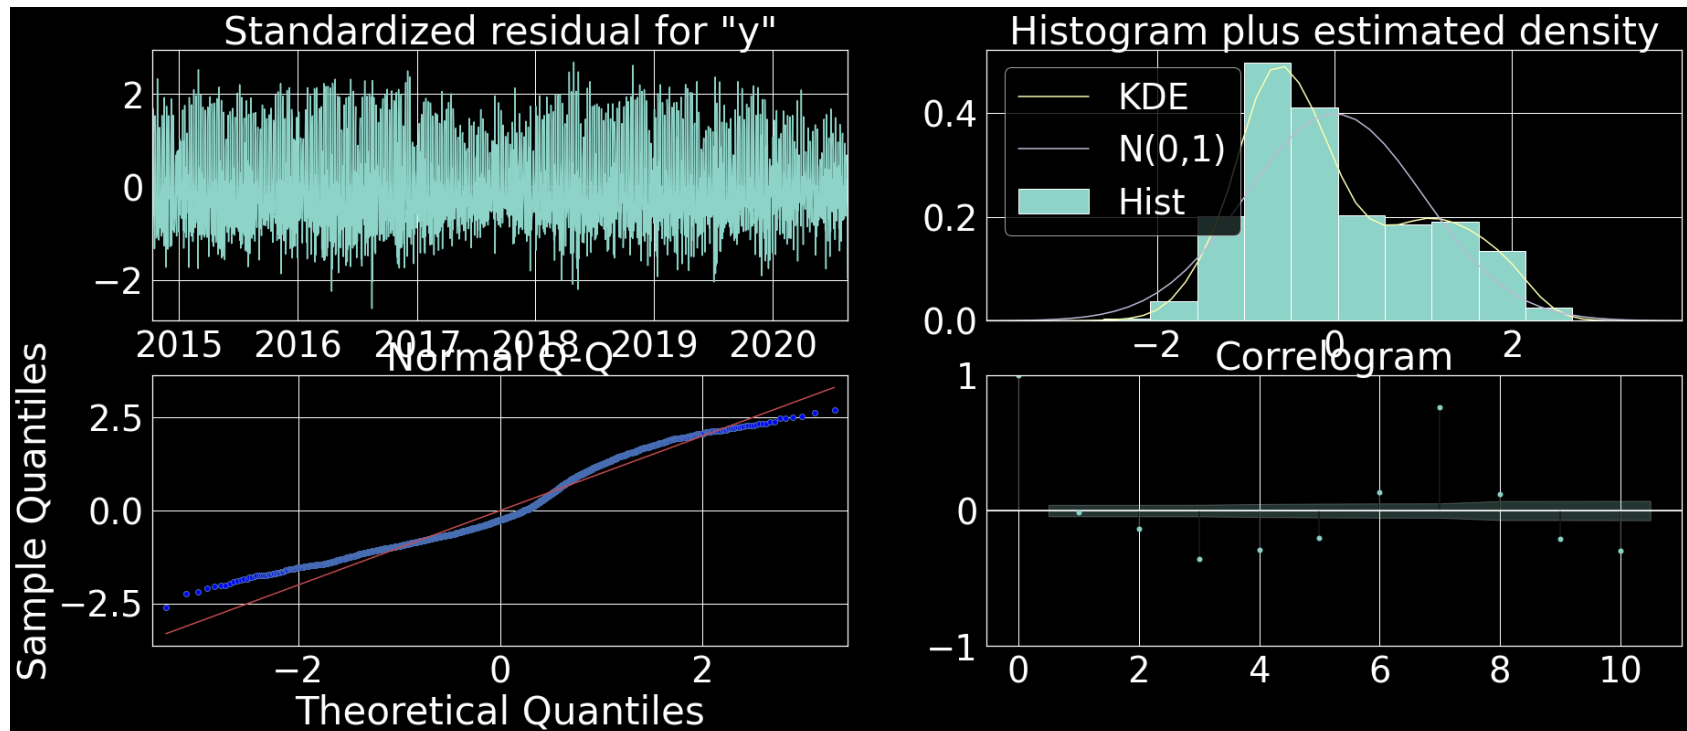
```
model = SARIMAX(data, exog=other_data, ...)
```

Examples of exogenous variables: Population, holidays, number of airline companies, major events

```
In [94]: # SARIMA example
from statsmodels.tsa.statespace.sarimax import SARIMAX

# fit model
model = SARIMAX(train["Page_Loads"], order=(1, 0, 1), seasonal_order=(0, 1, 1, 12))
model_fit = model.fit(dispatch=False)
```

```
In [95]: results.plot_diagnostics(figsize=(30, 12))
plt.show()
```



```
In [96]: print(train.shape, test.shape)
```

```
(42, 3) (30, 4)
```

```
In [97]: start_index = test.index.min()  
end_index = test.index.max()
```

```
#Predictions
```

```
predictions = model_fit.predict(start=start_index, end=end_index)
```

```
In [98]: from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, median_absolute_error, mean_squared_lo
```

```
In [99]: from math import sqrt
```

```
In [100... predictions.head(3)
```

```
Out[100]: 2018-03-31    128713.686417  
2018-04-30    130532.723844  
2018-05-31    124506.155304  
Freq: M, Name: predicted_mean, dtype: float64
```

```
In [101... test2 = test["Page_Loads"]
```

```
In [102... test2.head(3)
```

```
Out[102]: Date  
2018-03-31    151204  
2018-04-30    175572  
2018-05-31    161708  
Name: Page_Loads, dtype: int64
```

```
In [103... train2 = train["Page_Loads"]
```

```
In [104... train2.head(3)
```

```
Out[104]: Date  
2014-09-30     56052  
2014-10-31    121983  
2014-11-30    114190  
Name: Page_Loads, dtype: int64
```

```
In [105... # report performance
mse = mean_squared_error(test2[start_index:end_index], predictions)
rmse = sqrt(mse)
print('RMSE: {}, MSE:{}'.format(rmse,mse))
```

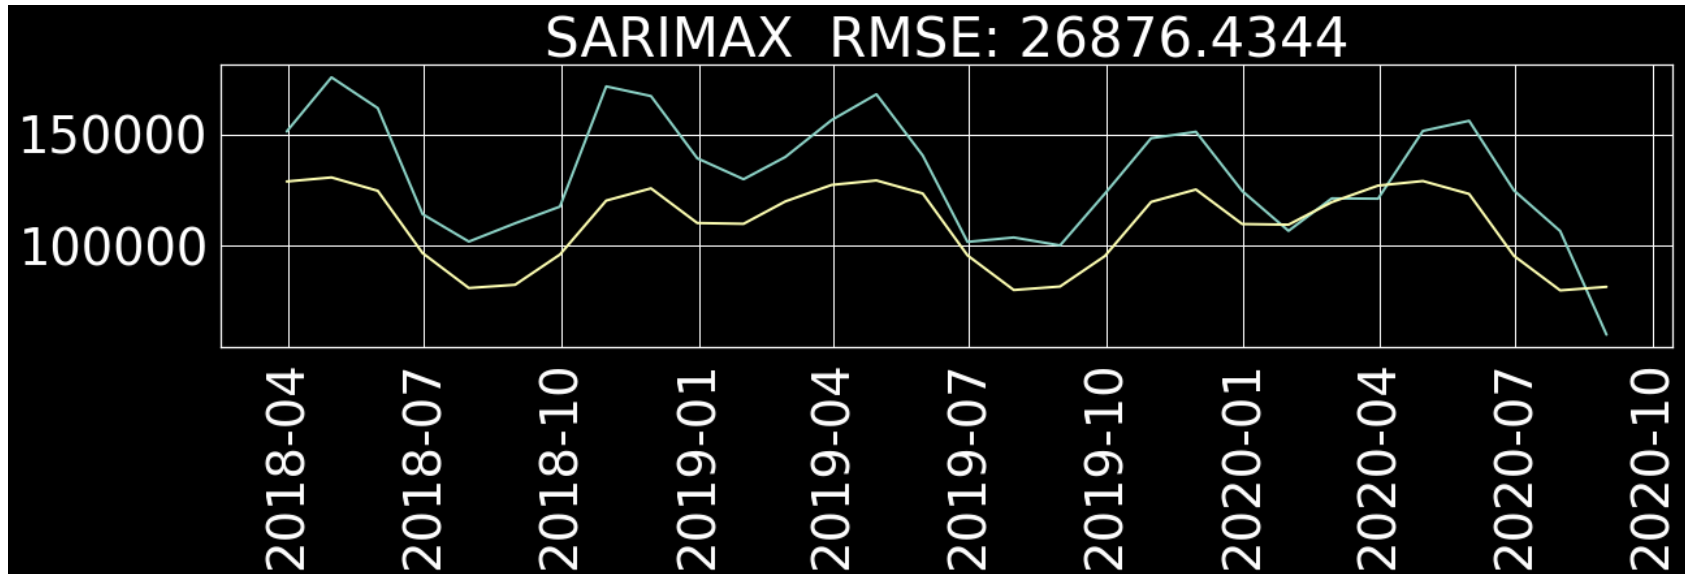
RMSE: 26876.43436509159, MSE:722342724.1810763

```
In [106... ###plt.xticks(x, labels, rotation='vertical')
```

```
In [110... import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (20,4)

plt.plot(test2, linewidth=2)
plt.plot(predictions, linewidth=2)
plt.xticks(rotation='vertical')
plt.title('SARIMAX RMSE: %.4f'% rmse)
```

Out[110]: Text(0.5, 1.0, 'SARIMAX RMSE: 26876.4344')



```
In [ ]: #####!pip install pmdarima
```

```
In [118... test.head(3)
```

Out[118]:

	Page_Loads	Page_Loads_First_Order_Differencing	Page_Loads_Second_Order_Differencing	Predicted_ARIMA
Date				
2018-03-31	151204	29464.0	19440.0	-22571.0
2018-04-30	175572	24368.0	-5096.0	-44621.0
2018-05-31	161708	-13864.0	-38232.0	-66671.0

In [119... `train.head(3)`

Out[119]:

	Page_Loads	Page_Loads_First_Order_Differencing	Page_Loads_Second_Order_Differencing
Date			
2014-09-30	56052	NaN	NaN
2014-10-31	121983	65931.0	NaN
2014-11-30	114190	-7793.0	-73724.0

In [113... `from pmdarima.arma import auto_arma`
`from pmdarima.arma import ADFTTest`

Hyper Parameter Tuning Using AUTO-ARIMA

In [120... `model=auto_arma(train["Page_Loads"],start_p=0,d=1,start_q=0,`
`max_p=5,max_d=5,max_q=5, start_P=0,`
`D=1, start_Q=0, max_P=5,max_D=5,`
`max_Q=5, m=12, seasonal=True,`
`error_action='warn',trace=True,`
`supress_warnings=True,stepwise=True,`
`random_state=20,n_fits=50)`

Performing stepwise search to minimize aic

```
ARIMA(0,1,0)(0,1,0)[12]      : AIC=632.135, Time=0.13 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=627.435, Time=0.27 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=623.573, Time=0.25 sec
ARIMA(0,1,1)(0,1,0)[12]      : AIC=634.064, Time=0.05 sec
ARIMA(0,1,1)(1,1,1)[12]      : AIC=623.847, Time=0.38 sec
ARIMA(0,1,1)(0,1,2)[12]      : AIC=624.012, Time=0.55 sec
ARIMA(0,1,1)(1,1,0)[12]      : AIC=628.068, Time=0.23 sec
ARIMA(0,1,1)(1,1,2)[12]      : AIC=625.842, Time=1.18 sec
ARIMA(0,1,0)(0,1,1)[12]      : AIC=623.990, Time=0.20 sec
ARIMA(1,1,1)(0,1,1)[12]      : AIC=624.205, Time=0.67 sec
ARIMA(0,1,2)(0,1,1)[12]      : AIC=625.096, Time=0.45 sec
ARIMA(1,1,0)(0,1,1)[12]      : AIC=623.597, Time=0.22 sec
ARIMA(1,1,2)(0,1,1)[12]      : AIC=625.334, Time=2.75 sec
ARIMA(0,1,1)(0,1,1)[12] intercept : AIC=624.010, Time=0.32 sec
```

Best model: ARIMA(0,1,1)(0,1,1)[12]

Total fit time: 7.716 seconds

In [121...

```
# SARIMA example
from statsmodels.tsa.statespace.sarimax import SARIMAX

# fit model
model_auto = SARIMAX(train["Page_Loads"], order=(0, 1, 1), seasonal_order=(0, 1, 1, 12))
model_auto_fit = model_auto.fit(dispatch=False)
```

In [122...

```
model_auto_fit.summary()
```


Out[122]:

SARIMAX Results

Dep. Variable:	Page_Loads		No. Observations:	42		
Model:	SARIMAX(0, 1, 1)x(0, 1, 1, 12)			Log Likelihood	-308.786	
Date:	Sun, 11 Dec 2022			AIC	623.573	
Time:	16:31:28			BIC	627.675	
Sample:	09-30-2014			HQIC	624.858	
	- 02-28-2018					
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
ma.L1	0.0660	0.213	0.310	0.757	-0.351	0.483
ma.S.L12	-0.6855	0.298	-2.302	0.021	-1.269	-0.102
sigma2	9.773e+07	1.18e-09	8.29e+16	0.000	9.77e+07	9.77e+07
Ljung-Box (L1) (Q):	0.21	Jarque-Bera (JB):	2.21			
Prob(Q):	0.65	Prob(JB):	0.33			
Heteroskedasticity (H):	1.11	Skew:	0.63			
Prob(H) (two-sided):	0.87	Kurtosis:	2.49			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 6.68e+32. Standard errors may be unstable.

In [123]...

```
print(train.shape, test.shape)
```

```
(42, 3) (30, 4)
```

```
In [125... start_index = test.index.min()
end_index = test.index.max()

#Predictions
## pred = model_auto_fit.get_prediction(start=start_index,end=end_index, dynamic=False)
```

```
In [127... print(start_index)
print(end_index)

2018-03-31 00:00:00
2020-08-31 00:00:00
```

```
In [128... predictions = model_auto_fit.predict(start=start_index, end=end_index)
```

```
In [132... print(predictions.shape, test.shape)

(30,) (30, 4)
```

```
In [133... test.columns
```

```
Out[133]: Index(['Page_Loads', 'Page_Loads_First_Order_Differencing',
        'Page_Loads_Second_Order_Differencing', 'Predicted_ARIMA'],
        dtype='object')
```

```
In [134... test3 = test["Page_Loads"]
```


```
In [135... # report performance
mse = mean_squared_error(test3[start_index:end_index], predictions)
rmse = sqrt(mse)
print('RMSE: {}, MSE:{}'.format(rmse,mse))

RMSE: 21696.720653403383, MSE:470747687.11182094
```

```
In [136... import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (20,4)

plt.plot(test3, linewidth=2)
plt.plot(predictions, linewidth=2)
plt.xticks(rotation='vertical')
plt.title('SARIMAX RMSE: %.4f'% rmse)
```

```
Out[136]: Text(0.5, 1.0, 'SARIMAX RMSE: 21696.7207')
```



In []: